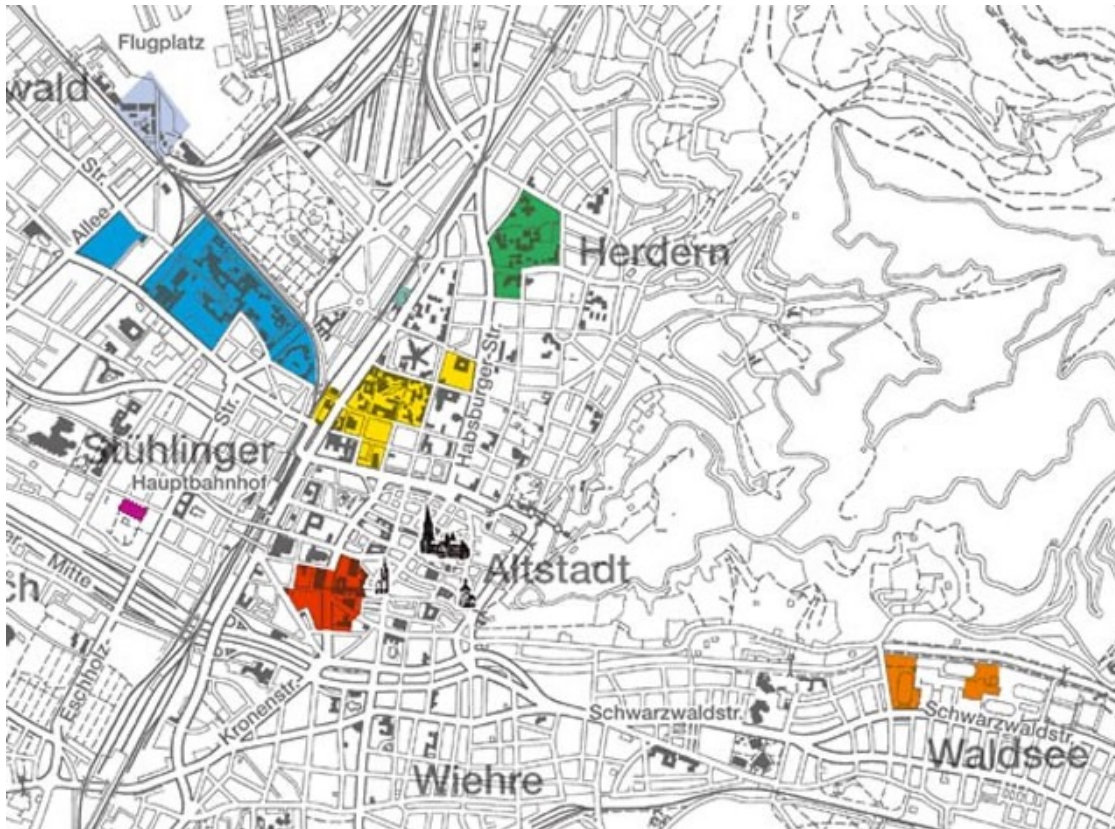


# R-Trees

# Outline

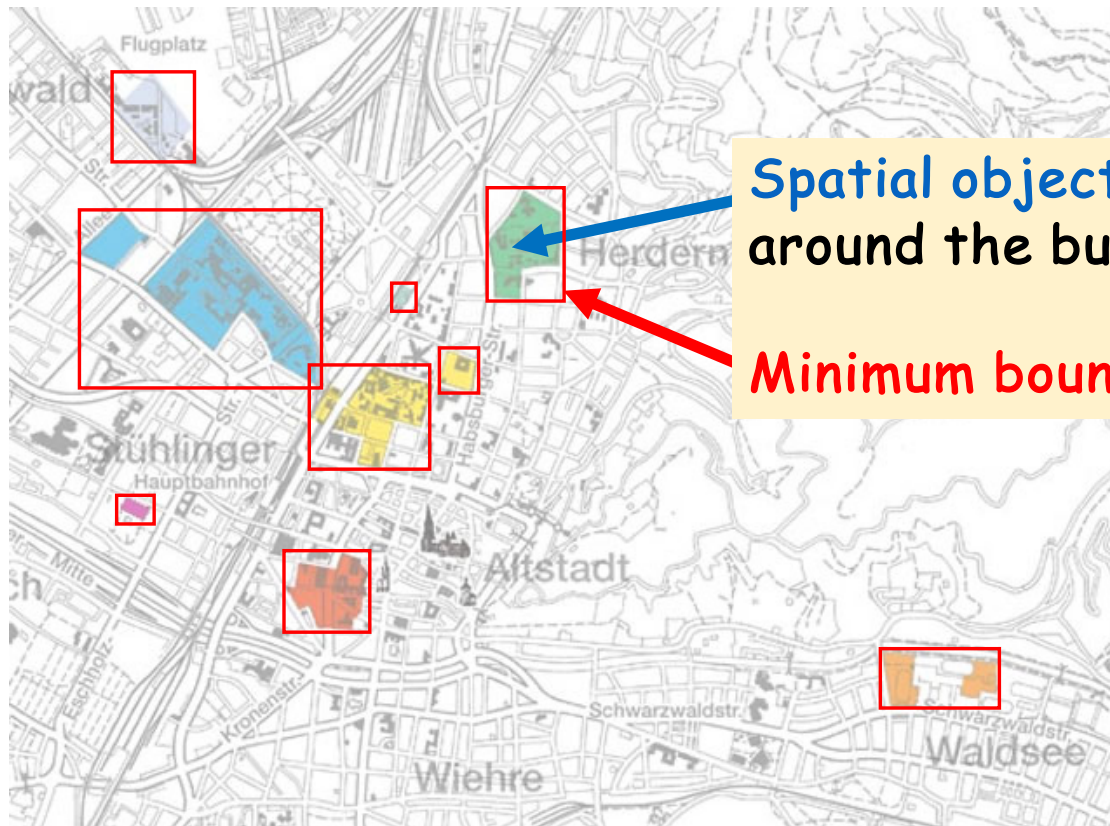
- Spatial data
- Shall have balanced-tree (B-tree) knowledge
- R-tree data structure
- Unique issues in R-tree (vs B-tree)
  - Set of continuous data items (vs a single data item)
  - Representing continuous data and optimizing their placements

# Spatial Database



Given a city map,  
"index" all university  
buildings in an  
efficient structure for  
quick topological  
search.

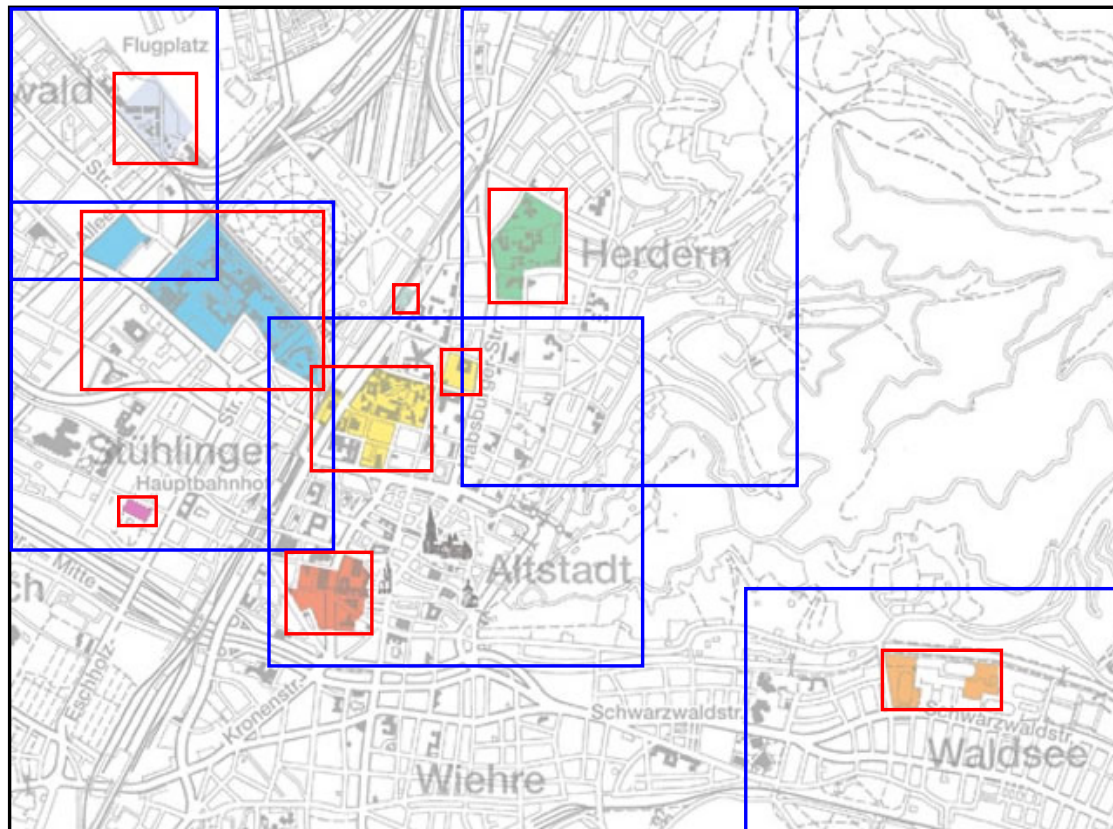
# Spatial Database (cont'd)



**Spatial object:** Contour (outline) of the area around the building(s).


**Minimum bounding region (MBR)** of the object.

# Spatial Database (cont'd)



# Recall: B-tree

- A B-Tree is an **ordered, dynamic, multi-way structure** of order  $m$  (i.e. each node has at most  $m$  children).
- The keys and the subtrees are arranged in the fashion of a **search tree**.

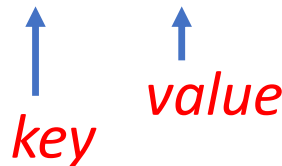


# Implementation in B-Tree

- **Insert**
  - Combine
  - Split
- **Delete**
  - Rotate
  - Combine
  - Split

# The R-Tree Index Structure

- An R-Tree is a **height-balanced** tree.
  - Leaves in the structure all appear **on the same level**.
- Index records in the leaf nodes contain **pointers** to the actual spatial-objects they represent. (akin to B+-tree)
- A spatial database consists of a collection of tuples representing spatial objects, known as **Entries**.
  - Each **Entry** has a unique **identifier** that **points** to one spatial object, and its MBR; i.e. Entry = (MBR, pointer).

  
*key* *value*




# R-Tree Index Structure - Leaf Entries



- An entry  $E$  in a *leaf* node is defined as:

$$E = (I, \text{tuple-identifier})$$

- $I$  (or *key*) refers to the *smallest bounding*  $n$ -dimensional region (MBR) that encompasses the spatial data pointed to by its *tuple-identifier* (or *value*).
- $I$  is a series of closed-intervals that make up each dimension of the binding region.



- Example. In 2D,  $I = (Ix, Iy)$ , where  $Ix = [xa, xb]$ , and  $Iy = [ya, yb]$ .





# R-Tree Index Structure - **Non-Leaf Entries**

- An **entry**  $E$  in a **non-leaf** node is defined as:

$$E = (I, \text{child-pointer})$$

where the *child-pointer* points to the child of this node, and  $I$  is the **MBR** that “encompasses” all the regions in the **child-node's pointer's entries**.



# Properties (B-tree, Essentially)

- Let  $M$  be the maximum number of entries that will fit in one node.
- Let  $m \leq M/2$  be a parameter specifying the minimum number of entries in one node.
  - Our previous lecture regarding B-tree lets  $m = \text{ceil}[M/2]$ .
- Then an R-Tree must satisfy the following properties:
  1. Every leaf node contains between  $m$  and  $M$  index records, unless it is the root.
  2. For each index-record Entry ( $I$ , *tuple-identifier*) in a leaf node,  $I$  is the MBR that spatially contains the  $n$ -dimensional data object represented by the *tuple-identifier*.
  3. Every non-leaf node has between  $m$  and  $M$  children, unless it is the root.
  4. For each Entry ( $I$ , *child-pointer*) in a non-leaf node,  $I$  is the MBR that spatially contains the regions in the child node.
  5. The root has two children unless it is a leaf.
  6. All leaves appear on the same level.

# Node Overflow and Underflow (Akin to B-tree)

- Node overflow



- A Node-Overflow happens when a new Entry is added to a fully packed node, causing the resulting number of entries in the node to exceed the upper-bound  $M$ .
- The 'overflow' node must be split, and all its current entries, as well as the new one, consolidated for local optimum arrangement.

- Node underflow

- A Node-Underflow happens when one or more Entries are removed from a node, causing the remaining number of entries in that node to fall below the lower-bound  $m$ .
- The underflow node must be condensed, and its entries dispersed for global optimum arrangement.

# Observations

- MBR's may overlap
  - in a **single node**
    - 1<sup>st</sup> level : (T, U) in node 1
    - 2<sup>nd</sup> level: (P, Q) in node 2
  - in **multiple nodes**
    - 2<sup>nd</sup> level: (Q, R) in nodes 2 and 3
    - 一parent node裡的MBR's有overlaps，distinct child nodes裡的MBR's就可能也會發生overlaps (overlaps是會傳遞到descendent nodes)



# B-Tree vs R-Tree

- B-tree:
  - A key can only appear in a "single" key segment
  - Key segments are "disjoint"
- R-tree:
  - A key object (MBR) may overlap "multiple" range segments, though the range segments may be disjoint
  - Range segments may be "overlapped", however

# Search Strategy in R-Tree

- Let  $Q$  be the query region.
- Let  $T$  be the root of the R-Tree.
- Search all entry-records whose regions overlaps  $Q$ .
- for (::
  - **Search sub-trees:**
    - If  $T$  is not leaf, then apply **Search** on every child-node entry  $E$  whose  $I$  overlaps  $Q$ .
  - **Search leaf nodes:**
    - If  $T$  is leaf, then check each entry  $E$  in the leaf and return  $E$  if  $E.I$  overlaps  $Q$ .

# Search Issue

- The search algorithm descends the tree from the root.
- More than one subtree under a node visited may need to be searched.
- *Cannot guarantee good worst-case performance.*
  - Countered by the algorithms during insertion, deletion, and update that maintain the tree in a form that allows the search algorithm to eliminate irrelevant regions of the indexed space.
  - So that only data near the search area need to be examined.
  - Emphasis is on the optimal placement of spatial objects with respect to the spatial location of other objects in the structure.





# Insertion in R-Tree

- Let  $E = (I, \text{tuple-identifier})$  be the new entry to be inserted.
- Let  $T$  be the root of the R-Tree.
  - **Locate a leaf**  $L$  starting from  $T$  to insert  $E$ .
  - **Add**  $E$  to  $L$ . If  $L$  is already full (overflow), **split**  $L$  into  $L$  and  $L'$ .
  - **Propagate** MBR changes (enlarged or reduced) **upwards**.
  - **Grow tree taller** if node split propagation causes  $T$  to split.


Insertion:

Which ways to Descend?




# Node Splitting Issue

Bad?



Good?



# Derivatives of the R-Tree

- May consider constraints:
  - **Overlap**: Total area contained within 2 or more MBRs (in a node).
  - **Coverage**: Total area of all the MBRs of all nodes (in a node).
- **Overlap**:
  - tighter/smaller MBRs -> more MBR records -> size of R tree increases (bad) and good for search, however
- **Coverage**:
  - Enlarged MBRs -> sparser in a MBR -> less MBR records -> search overhead increases as more MBR records visited

# Variants

- R+-tree
- R\*-tree
- ...